

Practicals to Convolutional Networks

Aufgabe P 2. *Convolutional Networks, due 14 December in class*

In this practical you will implement a convolutional neural network (CNN) for image classification in the CIFAR10 dataset containing $32 \times 32 \times 3$ images. The data import is provided in the stubs.

- (a) Define the graph of the CNN with following convolutional, pooling, fully-connected layers:
- i) Convolutional layer with filter size $[3, 3]$, no strides, padding 'same', 32 feature maps, and a bias. Apply the rectified linear activation function (`tf.nn.relu`).
 - ii) Convolutional layer with filter size $[3, 3]$, no strides, padding 'same', 32 feature maps, and a bias. Apply the rectified linear activation function (`tf.nn.relu`).
 - iii) Max pool layer with filter size of $[2, 2]$, strides of 2 for the width and height of the image, padding 'same'.
 - iv) Convolutional layer with filter size $[3, 3]$, no strides, padding 'same', 64 feature maps, and a bias. Apply the rectified linear activation function (`tf.nn.relu`).
 - v) Max pool layer with filter size of $[2, 2]$, strides of 2 for the width and height of the image, padding 'same'.
 - vi) Fully connected layer with output of 512 neurons. Input is the output of max-pool layer reshaped to a vector.
 - vii) Fully connected layer with output of 10 neurons. Input is the output previous fully connected layer.
 - viii) Use the `tf.losses.sparse_cross_entropy` as loss. Use the sparse version of cross-entropy since the labels are imported as numbers and not as one-hot-encoded vectors. Implement the accuracy for the classification.
 - ix) Train the model with `tf.train.AdamOptimizer` and try different learning rates for 10 epochs and 20 epochs. The stubs contain tensorboard loggings which you should turn in together with the accuracy in training and test.
- (b) Introduce a drop-out layer after the first fully connected layer with a keep probability of 0.5. Use your pretrained network and start a refinement training. Is the test accuracy better than the previous one?

Note: initialize the filter variables with `tf.contrib.layers.xavier_initializer_conv2d()` and the bias variables with `tf.constant_initializer(0.0001)`.